

1. Executable files

An executable file is a file containing instructions. These instructions must be in a format the CPU can understand. The standard format for the executable files on Linux is ELF¹ (*Executable and Linkable Format*). The format is defined in the `elf.h` header file. An executable file is a *binary* file, other binary files are *relocatable object files*, *core files*, *shared objects*.

1.1 The Executable and Linkable Format (ELF)

An ELF executable file consists of a an *ELF Header*, a *Program Header Table* or a *Section Header Table*, or both.

There are specific types defined for the ELF format which are listed in the following table²

Elf32_Addr	Unsigned program address, uint32_t
Elf32_Off	Unsigned file offset, uint32_t
Elf32_Section	Unsigned section index, uint16_t
Elf32_Versym	Unsigned version symbol information, uint16_t
Elf_Byte	unsigned char
Elf32_Half	uint16_t
Elf32_Sword	int32_t
Elf32_Word	uint32_t
Elf32_Sxword	int64_t
Elf32_Xword	uint64_t

Tab 1. ELF data types

For each type, both the 32-bit and 64-bit versions are defined. The same for the structures of the headers.

The *ELF Header* (*Ehdr*) is defined by the `Elf32_Ehdr` structure, which is defined as follows.

```
typedef struct
{
    unsigned char e_ident[EI_NIDENT]; /* Magic number and other info */
    Elf32_Half e_type; /* Object file type */
    Elf32_Half e_machine; /* Architecture */
    Elf32_Word e_version; /* Object file version */
    Elf32_Addr e_entry; /* Entry point virtual address */
    Elf32_Off e_phoff; /* Program header table file offset */
    Elf32_Off e_shoff; /* Section header table file offset */
    Elf32_Word e_flags; /* Processor-specific flags */
    Elf32_Half e_ehsize; /* ELF header size in bytes */
}
```

¹ See the `elf(5)` manual page

² In the `elf.h` header are also defined the 64-bit version of the data types

```

Elf32_Half e_phentsize;    /* Program header table entry size */
Elf32_Half e_phnum;        /* Program header table entry count */
Elf32_Half e_shentsize;    /* Section header table entry size */
Elf32_Half e_shnum;        /* Section header table entry count */
Elf32_Half e_shstrndx;     /* Section header string table index */
} Elf32_Ehdr;

```

The 64-bit version of this structure is `Elf64_Ehdr`.

The *Program Header Table* (Phdr) is an array of structures found in *executable* and *shared object files*. These structures describe the segments and other information required by the system to load and prepare the program for execution. A *segment* in an *object file* may contain one or more sections.

The two fields `e_phentsize` and `e_phnum` of the `Elf32_Ehdr` structure specify the *program header size*.

The `Elf32_Phdr` structure is defined as follows.

```

typedef struct
{
    Elf32_Word p_type;        /* Segment type */
    Elf32_Off p_offset;       /* Segment file offset */
    Elf32_Addr p_vaddr;       /* Segment virtual address */
    Elf32_Addr p_paddr;       /* Segment physical address */
    Elf32_Word p_filesz;      /* Segment size in file */
    Elf32_Word p_memsz;       /* Segment size in memory */
    Elf32_Word p_flags;       /* Segment flags */
    Elf32_Word p_align;       /* Segment alignment */
} Elf32_Phdr;

```

The 64-bit version of the structure is `Elf64_Phdr`.

The *Section Header Table* (Shdr) is an array of `Elf32_Shdr` (or `Elf64_Shdr`) structures that allows to locate all the file's sections.

The `e_shoff` field of the ELF Header specifies the byte offset from the beginning of the file to the *Section Header Table* (Shdr).

The `e_shnum` field of the ELF Header (Ehdr) specifies the number of entries contained in the *Section Header Table*, while the `e_shentsize` field specifies the size, in bytes, of each entry.

Some array indices are reserved to represent special positions or predefined meanings.³

The `Elf32_Shdr` structure is defined as follows

```

typedef struct
{
    Elf32_Word sh_name;       /* Section name (string tbl index) */
    Elf32_Word sh_type;       /* Section type */
    Elf32_Word sh_flags;      /* Section flags */

```

3 These indexes are listed in the `elf(5)` manual page

```

Elf32_Addr sh_addr;      /* Section virtual addr at execution */
Elf32_Off  sh_offset;    /* Section file offset */
Elf32_Word sh_size;      /* Section size in bytes */
Elf32_Word sh_link;      /* Link to another section */
Elf32_Word sh_info;      /* Additional section information */
Elf32_Word sh_addralign; /* Section alignment */
Elf32_Word sh_entsize;   /* Entry size if section holds table */
} Elf32_Shdr;

```

The 64-bit version of the structure is `Elf64_Shdr`.

In addition to the three main structures, ELF files also include the *Symbol Table* (Sym), the *Relocation Entries* (Rel and Rela), and the *Dynamic Tags* (Dyn).

The *object files Symbol Table* uses strings to represent *symbol* and *section names*. This table contains information to locate and relocate the symbolic definitions and references of a program, and is defined as follows.

```

typedef struct
{
    Elf32_Word st_name;      /* Symbol name (string tbl index) */
    Elf32_Addr st_value;     /* Symbol value */
    Elf32_Word st_size;     /* Symbol size */
    unsigned char st_info;   /* Symbol type and binding */
    unsigned char st_other;  /* Symbol visibility */
    Elf32_Section st_shndx; /* Section index */
} Elf32_Sym;

```

The 64-bit version of the structure is `Elf64_Sym`.

Some symbols can have additional information that are stored in the `Elf32_Syminfo` structure, which is defined as follows.

```

typedef struct
{
    Elf32_Half si_boundto; /* Direct bindings, symbol bound to */
    Elf32_Half si_flags;   /* Per symbol flags */
} Elf32_Syminfo;

```

The 64-bit version of the structure is `Elf64_Syminfo`.

The *Relocation Entries* are defined by the two structures `Elf32_Rel` and `Elf32_Rela`, where `Elf32_Rela` defines an entry that need an addend.

Relocation means to connect a *symbolic reference* to a *symbolic definition*. The *relocatable files* must have information that describes how to modify their section contents in order to allow the *executable* and *shared object files* to know the right information for a process's *program image*. A *process's program image* is the program loaded into the memory. Once loaded it must know where to find the symbols it uses, this information is stored in the *relocation entries table*.

The `Elf32_Rel` structure is defined as follows.

```
typedef struct
{
    Elf32_Addr r_offset;      /* Address */
    Elf32_Word r_info;       /* Relocation type and symbol index */
} Elf32_Rel;
```

The 64-bit version of the structure is `Elf64_Rel`.

The `Elf32_Rela` structure is defined as follows.

```
typedef struct
{
    Elf32_Addr r_offset;      /* Address */
    Elf32_Word r_info;       /* Relocation type and symbol index */
    Elf32_Sword r_addend;     /* Addend */
} Elf32_Rela;
```

The 64-bit version of the structure is `Elf64_Rela`.

The *Dynamic Tags* (Dyn) contains structures holding information about the *dynamic linking*. The `Elf32_Dyn` structure is defined as follows.

```
typedef struct
{
    Elf32_Sword d_tag;        /* Dynamic entry type */
    union
    {
        Elf32_Word d_val;     /* Integer value */
        Elf32_Addr d_ptr;     /* Address value */
    } d_un;
} Elf32_Dyn;
```

The 64-bit version of the structure is `Elf64_Dyn`.

Another structure used by ELF is the *Notes Header*. This contains arbitrary information used by the system. This data is largely used by *core files*. Also the *GNU tool chain* uses this header to pass information from the *linker* to the *C library*.

The `Elf32_Nhdr` structure is defined as follows.

```
typedef struct
{
    Elf32_Word n_namesz;      /* Length of the note's name. */
    Elf32_Word n_descsz;      /* Length of the note's descriptor. */
    Elf32_Word n_type;        /* Type of the note. */
} Elf32_Nhdr;
```

The 64-bit version of the structure is `Elf64_Nhdr`.

The `elf.h` header contains many other structures to specify all the features of the ELF format. The aforementioned are the most commonly used. The following list summarizes the structures we have seen so far.

Ehdr (ELF Header)

Contains general information about the ELF file, such as file type, architecture, entry point address, and offsets to other structures

Phdr (Program Header Table)

Describes the segments used by the system loader to prepare the program for execution

Shdr (Section Header Table)

Describes the sections contained in the ELF file, such as `.text`, `.data`, and `.bss`

Sym (Symbol Table)

Stores information about symbols, including functions and global variables

Syminfo (Symbol Information Table)

Provides additional information about symbols, mainly used for dynamic linking and symbol resolution

Rel / Rela (Relocation Entries)

Contain information required to modify addresses and references during linking or loading

Dyn (Dynamic Section/Tags)

Contains information needed for dynamic linking, such as shared libraries and runtime linking data

Note (Notes Header/Section)

Stores auxiliary information such as ABI details, build IDs, or vendor-specific metadata

1.2 The ELF Header

On Linux systems, there are many available tools to analyze *binary files*: `hexedit`, `objdump`, `gdb`, `readelf`, etc.. The `readelf` program reads and display the information contained in the ELF file passed as argument. We take a look at the *ELF Header* of the following program:

```
#include <stdio.h>
int main() {
    printf("Hello Linux World!!");
    return 0;
}
```

```
[linux@memory] readelf -h hello
```

```
ELF Header:
```

```
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                                ELF64
  Data:                                      2's complement, little endian
  Version:                                1 (current)
  OS/ABI:                                  UNIX - System V
  ABI Version:                            0
  Type:                                    DYN (Position-Independent Executable file)
  Machine:                                Advanced Micro Devices X86-64
  Version:                                0x1
  Entry point address:                    0x1050
  Start of program headers:               64 (bytes into file)
  Start of section headers:              13968 (bytes into file)
  Flags:                                  0x0
  Size of this header:                    64 (bytes)
  Size of program headers:                56 (bytes)
  Number of program headers:              14
  Size of section headers:                64 (bytes)
  Number of section headers:              31
  Section header string table index:      30
```

The output displays the following fields:

Magic

Identifies the file as ELF. Always starts with `0x7f 45 4c 46 (\x7fELF)`

Class

Whether the binary is 32-bit or 64-bit (ELF32 or ELF64)

Data

Endianness, little-endian or big-endian

Version

ELF format version (1 for current ELF)

OS/ABI

Target ABI (e.g. UNIX System V, Linux, FreeBSD), usually "UNIX - System V" for Linux

ABI Version

Version of the ABI⁴

Type

The type of the ELF file

ET_EXEC → Executable

ET_REL → Relocatable object file (.o)

ET_DYN → Shared object (.so, PIE executable file)

ET_CORE → Core dump

Machine

The target architecture (x86-64, ARM, RISC-V, etc.)

⁴ ABI stands for *Application Binary Interface* and defines the low-level interface between the binary code and the operating system, see the links in the *Bibliography* for further details

Version

Version in hexadecimal format (0x1)

Entry point address

The address where the execution begins (0x1050)

Start of program headers

The file offset where the *program header table* begins, used by the kernel loader,

Start of section headers

The file offset where the *section header table* begins, used by tools like `ld`

Flags

Machine-specific flags, 0x0 for x86-64

Size of this header

64 for ELF64, 52 for ELF32

Size of program headers

Size of one entry in the *program header table*

Number of program headers

The number of the entries (*segments*)

Size of section headers

The size of one entry in the *section header table*

Number of section headers

The number of entries (*sections*)

Section header string table index

Index of the section header that contains the section names

1.3 Sections and Segments

Passing the option `-S` we can retrieve the *Section Headers*⁵.

```
[linux@memory] readelf -S --wide hello
```

```
There are 31 section headers, starting at offset 0x3690:
```

```
Section Headers:
```

[Nr]	Name	Type	Address	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	0000000000000000	000000	000000	00		0	0	0
[1]	.note.gnu.property	NOTE	0000000000000350	000350	000020	00	A	0	0	8
[2]	.note.gnu.build-id	NOTE	0000000000000370	000370	000024	00	A	0	0	4
[3]	.interp	PROGBITS	0000000000000394	000394	00001c	00	A	0	0	1
[4]	.gnu.hash	GNU_HASH	00000000000003b0	0003b0	000024	00	A	5	0	8
[5]	.dynsym	DYNSYM	00000000000003d8	0003d8	0000a8	18	A	6	1	8
[6]	.dynstr	STRTAB	0000000000000480	000480	00008f	00	A	0	0	1
[7]	.gnu.version	VERSYM	0000000000000510	000510	00000e	02	A	5	0	2
[8]	.gnu.version_r	VERNEED	0000000000000520	000520	000030	00	A	6	1	8
[9]	.rela.dyn	RELA	0000000000000550	000550	0000c0	18	A	5	0	8

5 The `--wide` option ensures that the output fits on one line without wrapping to the next line

[10]	.rela.plt	RELA	00000000000000610	000610	000018	18	AI	5	24	8
[11]	.init	PROGBITS	0000000000001000	001000	000017	00	AX	0	0	4
[12]	.plt	PROGBITS	0000000000001020	001020	000020	10	AX	0	0	16
[13]	.plt.got	PROGBITS	0000000000001040	001040	000008	08	AX	0	0	8
[14]	.text	PROGBITS	0000000000001050	001050	000108	00	AX	0	0	16
[15]	.fini	PROGBITS	0000000000001158	001158	000009	00	AX	0	0	4
[16]	.rodata	PROGBITS	0000000000002000	002000	000018	00	A	0	0	4
[17]	.eh_frame_hdr	PROGBITS	0000000000002018	002018	00002c	00	A	0	0	4
[18]	.eh_frame	PROGBITS	0000000000002048	002048	0000ac	00	A	0	0	8
[19]	.note.ABI-tag	NOTE	00000000000020f4	0020f4	000020	00	A	0	0	4
[20]	.init_array	INIT_ARRAY	0000000000003dd0	002dd0	000008	08	WA	0	0	8
[21]	.fini_array	FINI_ARRAY	0000000000003dd8	002dd8	000008	08	WA	0	0	8
[22]	.dynamic	DYNAMIC	0000000000003de0	002de0	0001e0	10	WA	6	0	8
[23]	.got	PROGBITS	0000000000003fc0	002fc0	000028	08	WA	0	0	8
[24]	.got.plt	PROGBITS	0000000000003fe8	002fe8	000020	08	WA	0	0	8
[25]	.data	PROGBITS	0000000000004008	003008	000010	00	WA	0	0	8
[26]	.bss	NOBITS	0000000000004018	003018	000008	00	WA	0	0	1
[27]	.comment	PROGBITS	0000000000000000	003018	00001e	01	MS	0	0	1
[28]	.symtab	SYMTAB	0000000000000000	003038	000360	18		29	18	8
[29]	.strtab	STRTAB	0000000000000000	003398	0001dd	00		0	0	1
[30]	.shstrtab	STRTAB	0000000000000000	003575	00011a	00		0	0	1

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
D (mbind), l (large), p (processor specific)

The output is composed by the columns:

[Nr]

The section index number (0 = reserved, NULL section)

Name

The section name (.text, .data, .bss, .symtab, .strtab). the names come from the *section header string table*

Type

The section type⁶

PROGBITS	→ raw program data (code or constants)
NOBITS	→ occupies memory but no space in file (. bss)
SYMTAB/DYNSYM	→ symbol tables
STRTAB	→ string tables
NOTE	→ notes
REL/RELA	→ relocation entries
NULL	→ unused

Address

The *virtual memory address* where the section is loaded at runtime

Off

The file offset (in bytes) where this section begins inside the ELF file

Size

The size of the section

⁶ The complete list is in the `elf.h` header file

ES (entry size)

Size of each entry if the section holds a table of fixed-size entries, if it is not a table 0

Flg (flags)

Section attributes⁷

A → Alloc (section is loaded into memory)

X → Executable (contains code)

W → Writable

M → Mergeable

S → Strings (null-terminated)

I → Info link used

Lk (link)

Section index link

For SYMTAB → points to the string table section index (.strtab)

For REL/RELA → index of the section with symbols used for relocation

Inf (info)

Extra information

For relocation section → index of the section to which relocations apply

For SYMTAB → number of the *section header table* index of the associated section

Al (align)

Required alignment of section in memory (power of 2), e.g. 16 → address must be multiple of 16

To retrieve the *Program Headers* we can use the `-l` option.

```
[linux@memory] readelf -l --wide hello
Elf file type is DYN (Position-Independent Executable file)
Entry point 0x1050
There are 14 program headers, starting at offset 64

Program Headers:
  Type           Offset    VirtAddr           PhysAddr           FileSiz  MemSiz   Flg
  Align
  PHDR           0x000040  0x0000000000000040 0x0000000000000040 0x000310 0x000310 R
  0x8
  INTERP         0x000394  0x0000000000000394 0x0000000000000394 0x00001c 0x00001c R
  0x1
    [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
  LOAD           0x000000  0x0000000000000000 0x0000000000000000 0x000628 0x000628 R
  0x1000
  LOAD           0x001000  0x0000000000000100 0x0000000000000100 0x000161 0x000161 R E
  0x1000
  LOAD           0x002000  0x0000000000000200 0x0000000000000200 0x000114 0x000114 R
  0x1000
  LOAD           0x002dd0  0x00000000000003dd 0x00000000000003dd 0x000248 0x000250 RW
  0x1000
  DYNAMIC        0x002de0  0x00000000000003de 0x00000000000003de 0x0001e0 0x0001e0 RW
  0x8
  NOTE           0x000350  0x0000000000000350 0x0000000000000350 0x000020 0x000020 R
  0x8
  NOTE           0x000370  0x0000000000000370 0x0000000000000370 0x000024 0x000024 R
  0x4
```

⁷ The complete list is in the `elf.h` header file

```

NOTE          0x0020f4 0x000000000000020f4 0x000000000000020f4 0x000020 0x000020 R
0x4
GNU_PROPERTY  0x000350 0x0000000000000350 0x0000000000000350 0x000020 0x000020 R
0x8
GNU_EH_FRAME  0x002018 0x00000000000002018 0x00000000000002018 0x00002c 0x00002c R
0x4
GNU_STACK     0x000000 0x0000000000000000 0x0000000000000000 0x000000 0x000000 RW
0x10
GNU_RELRO     0x002dd0 0x00000000000003dd0 0x00000000000003dd0 0x000230 0x000230 R
0x1

```

Section to Segment mapping:

Segment Sections...

```

00
01      .interp
02      .note.gnu.property .note.gnu.buildid .interp .gnu.hash .dynsym .dynstr
      .gnu.version .gnu.version_r .rela.dyn .rela.plt
03      .init .plt .plt.got .text .fini
04      .rodata .eh_frame_hdr .eh_frame .note.ABI-tag
05      .init_array .fini_array .dynamic .got .got.plt .data .bss
06      .dynamic
07      .note.gnu.property
08      .note.gnu.build-id
09      .note.ABI-tag
10      .note.gnu.property
11      .eh_frame_hdr
12
13      .init_array .fini_array .dynamic .got

```

The output is composed by the following columns:

Type

The segment type that tells the loader what to do

```

LOAD          → Load this segment into memory
DYNAMIC       → Dynamic linking info
INTERP       → Path to dynamic linker (ld.so)
NOTE         → Auxiliary info (build-id, ABI tags)
PHDR         → Location of the program header table itself
GNU_STACK    → Stack permissions (RW vs RWX)
GNU_RELRO    → Read-only after relocation

```

Offset

File offset of the segment in the ELF file

VirtAddr

Virtual memory address where the segment will be mapped at runtime

PhysAddr

Physical address (unused in Linux, matches VirtAddr)

FileSiz

Number of bytes to read from the file

MemSiz

Number of bytes to allocate in memory, if `MemSiz > FileSiz` the extra space is zero-filled⁸

Flags

Memory permissions for the segment

R → Read

W → Write

E → Execute

Align

Alignment requirement

The second part of the output shows the mapping between the *Segments* and the *Sections*.

The *Section Headers* are used by the *linker* and *debuggers* while the *Program headers* are used by the *kernel loader* at runtime to map the binary into memory.

The *Sections* are: `.text`, `.data`, `.rodata`, `.bss`, `.symtab`, `.strtab`. Each *Section* has a specific type (e.g. code, data, symbol table) and associated flags that define its properties, such as read, write, or execute permissions. Sections are primarily used by the *linker* during the process of building the executable.

<code>.text</code>	Compiled code
<code>.data</code>	Initialized data
<code>.rodata</code>	Constants, strings
<code>.bss</code>	Uninitialized data
<code>.rel.text</code> <code>.rela.text</code>	Relocation info
<code>.symtab</code>	Symbol table
<code>.strtab</code>	String table

Tab 2. Sections of an executable

Sections are created in object files (`.o`) during the compilation process. Once the object files contain these sections, the linker (`ld`) merges sections of the same type from all object files to produce the final executable or shared library.

The *linker* also resolves the *symbols*, applies the relocations and discards the `.symtab` if not building with debug info (`-g` option of `gcc`).

After the linking process, *sections* are no longer needed by the kernel. Instead, they are mainly used by tools such as debuggers, disassemblers, and other analysis utilities.

Once the linking process is completed, an executable file is produced that can be loaded into memory by the *kernel loader*.

⁸ That is how the `.bss` section is created

The *kernel loader* reads the *Program Header Table* (segments) to map the segments. *Segments* are contiguous memory regions containing several *Sections* together and have permissions like read, write and execute.

PIE (*Position Independent Executable*) files are mapped into memory by treating the addresses specified in the *Program Header Table* as offsets from a randomized base address chosen by the kernel.

We compile the following code

```
#include <stdio.h>

#define MESSAGE "Hello Linux World!!"

int main() {
    printf(MESSAGE);
    return 0;
}
```

passing the `-S` option to `gcc`, in order to stop the compilation before assembling⁹, and the option `-masm=intel` to use the Intel syntax for the code.

```
[linux@memory] gcc -masm=intel -S hello.c
```

The `hello.s` file is created.

```
.file "hello.c"
.intel_syntax noprefix
.text
.section .rodata
.LC0:
.string "Hello Linux World!!"
.text
.globl main
.type main, @function
main:
.LFB0:
.cfi_startproc
push rbp
.cfi_def_cfa_offset 16
.cfi_offset 6, -16
mov rbp, rsp
.cfi_def_cfa_register 6
lea rax, .LC0[rip]
mov rdi, rax
mov eax, 0
call printf@PLT
mov eax, 0
```

⁹ In this way we can create and analyze the assembly file

```

    pop rbp
    .cfi_def_cfa 7, 8
    ret
    .cfi_endproc
.LFE0:
    .size main, .-main
    .ident "GCC: (Debian 14.3.0-5) 14.3.0"
    .section .note.GNU-stack,"",@progbits

```

hello.s

We can see the *Sections* created by the compiler.

The `file` command executed on the `hello` executable give us information about the binary file.

```
[linux@memory] file hello
hello: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV), dynamically linked, interpreter
/lib64/ld-linux-x86-64.so.2, BuildID[sha1]=8bdc9b7b2228d196e34fe38dff9e9109ca5738d7, for
GNU/Linux 3.2.0, not stripped
```

The output shows that the program is an ELF 64-bit little-endian (LSB) PIE executable compiled for the x86-64 architecture. The binary is *dynamically linked*, meaning it depends on *shared libraries* at runtime, and it specifies the *dynamic linker* `/lib64/ld-linux-x86-64.so.2` used to load the program.

The output also includes a BuildID, which is a unique hash used to identify the binary version. It indicates compatibility with GNU/Linux 3.2.0 or later. Finally, the program is not stripped, meaning it still contains debugging symbols and metadata that can be used for analysis.

LSB	→ Least Significant Byte first (<i>little-endian</i>)
PIE	→ Position Independent Executable
not stripped	→ the ELF file still contains (after the compilation) the <i>section headers</i> (<code>.symtab</code> , <code>.strtab</code> , <code>.debug_*</code>) that are not required for the execution

To remove the *section headers* from the executable we can use the `strip`¹⁰ command

```
[linux@memory] strip hello

[linux@memory] file hello
hello: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV), dynamically linked, interpreter
/lib64/ld-linux-x86-64.so.2, BuildID[sha1]=8bdc9b7b2228d196e34fe38dff9e9109ca5738d7, for
GNU/Linux 3.2.0, stripped
```

After running `strip`, the executable is marked as *stripped*, meaning that the section headers and most debugging-related information have been removed from the binary. As shown by the `file` output, the program remains fully functional as an ELF executable, but it no longer contains symbols and metadata useful for analysis.

¹⁰ See the `strip(1)` manual page

Section headers are mainly useful for development and analysis tools such as debuggers, disassemblers, and reverse engineering utilities. They are not required by the kernel at runtime, since the kernel relies on the *program headers* instead to load and map the executable into memory.

We can also see a difference between the size of the *not stripped* and *stripped* file

```
[linux@memory] ls -l hello  
-rwxrwxr-x 1 io io 15952 Sep 30 18:18 hello
```

```
[linux@memory] strip hello
```

```
[linux@memory] ls -l hello  
-rwxrwxr-x 1 io io 14464 Sep 30 18:19 hello
```